



μ TVM: Running the TVM Stack on Bare Metal

TVM Conf 2020

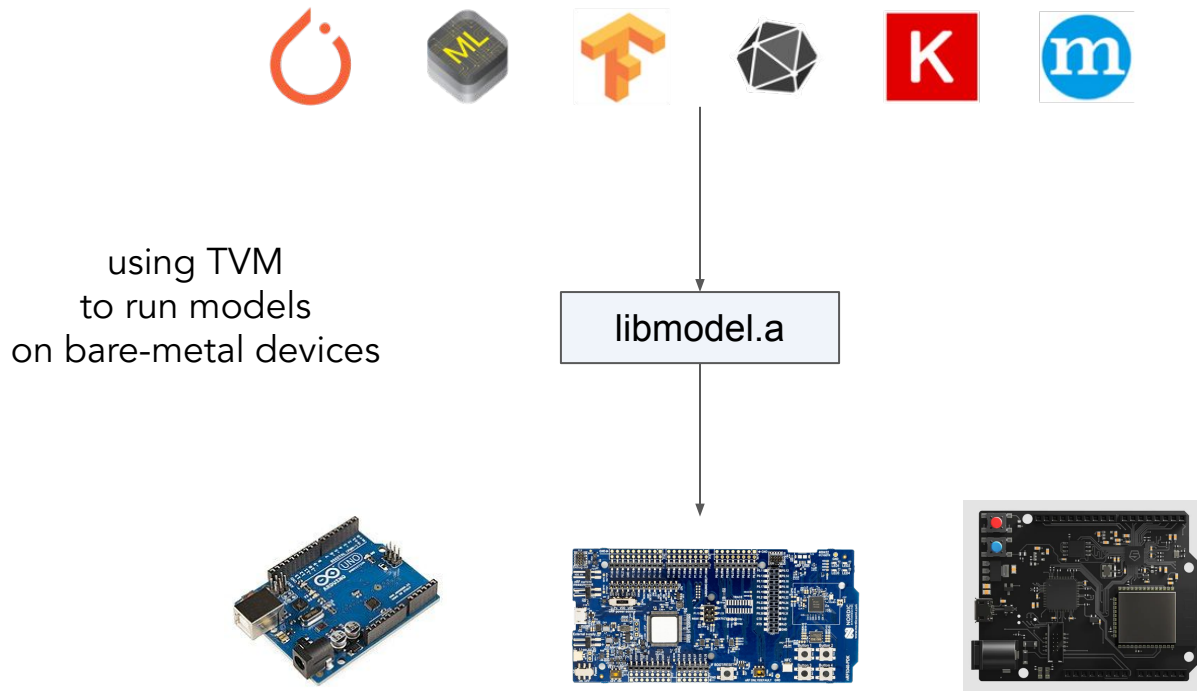
Andrew Reusch

Outline

- What is μ TVM?
- How μ TVM Works
- Demo Walkthrough
- Future Directions
- Q&A



What is μ TVM?



...Bare Metal?

To μ TVM, bare metal is not just:



Raspberry Pi

- These (usually) have operating systems



Reserved Cloud Instances

- These still have Virtual Memory



Running outside a VM


- Traditional TVM can run here




...Bare Metal?

Bare metal is (often) IoT-class devices

- ⚡ AzureSphere
- ⚡ Arduino
- ⚡ Cortex-M class micro-controllers



μ TVM works in places without...




Operating Systems

- no files, DLLs, .so, memory mapping, kernels



Virtual Memory

- No malloc, C++ RAII, exceptions, ...





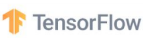



Advanced Programming Languages

- No C++, Rust, Python, ...
(But we like those and you could use them!)



The deployment challenge



 PyTorch	?	?	?	?	?	?	?	?	?	?
 Caffe2	?	?	?	?	?	?	?	?	?	?
 TensorFlow	?	?	?	?	?	?	?	?	?	?
 mxnet	?	?	?	?	?	?	?	?	?	?
 ONNX	?	?	?	?	?	?	?	?	?	?
 TensorFlow Lite	?	?	?	?	?	?	?	?	?	?

The deployment challenge

frameworks

↓

										
	?	?	?	?	?	?	?	?	?	?
	?	?	?	?	?	?	?	?	?	?
	?	?	?	?	?	?	?	?	?	?
	?	?	?	?	?	?	?	?	?	?
	?	?	?	?	?	?	?	?	?	?
	?	?	?	?	?	?	?	?	?	?

The deployment challenge

targets →

										
 PyTorch	?	?	?	?	?	?	?	?	?	?
 Caffe2	?	?	?	?	?	?	?	?	?	?
 TensorFlow	?	?	?	?	?	?	?	?	?	?
 mxnet	?	?	?	?	?	?	?	?	?	?
 ONNX	?	?	?	?	?	?	?	?	?	?
 TensorFlow Lite	?	?	?	?	?	?	?	?	?	?







The deployment challenge

targets →

										
 PyTorch	✓	?	?	✓	?	?	?	?	?	?
 Caffe2	✓	?	?	✓	?	?	?	?	?	?
 TensorFlow	✓	?	?	✓	?	?	?	?	?	?
 mxnet	✓	?	?	✓	?	?	?	?	?	?
 ONNX	✓	?	?	✓	?	?	?	?	?	?
 TensorFlow Lite	✓	?	?	✓	?	?	?	?	?	?







The deployment challenge



 PyTorch	?	?	?	?	?	?	?	?	?	?
 Caffe2	?	?	?	?	?	?	?	?	?	?
 TensorFlow	?	?	?	?	?	?	?	?	?	?
 mxnet	?	?	?	?	?	?	?	?	?	?
 ONNX	?	?	?	?	?	?	?	?	?	?
 TensorFlow Lite	?	✓	?	?	?	?	?	?	?	?

The deployment challenge



 PyTorch	?	?	?	?	?	?	?	?	?	?
 Caffe2	?	✗	?	?	?	?	?	?	?	?
 TensorFlow	?	?	?	?	?	?	?	?	?	?
 mxnet	?	?	?	?	?	?	?	?	?	?
 ONNX	?	?	?	?	?	?	?	?	?	?
 TensorFlow Lite	?	?	?	?	?	?	?	?	?	?

Bare Metal Deployment Challenges

- Less abstraction than full OS
 - Less tools to work with
- Resources are tighter
 - Scheduling is harder
- Demands are unique per-chip and per-project
 - Code reuse is tricky



The μ TVM Approach

- Batteries Included
 - μ TVM can be used with only the standard C library
- Compute-centric
 - μ TVM does not configure the SoC--it only runs computations
 - μ TVM integrates with RTOS like Zephyr and mBED for SoC configuration
- Transparent
 - μ TVM binaries can be compiled directly from source

How μ TVM Works



```
int32_t fused_conv2d_right_shift_add() {  
    // ...  
}
```

How μ TVM Works




tvm

```
int main() {  
    // configure SoC  
    TVMInitializeRuntime();  
    TVMGraphRuntime_Run();  
}
```

+

```
int32_t fused_conv2d_right_shift_add() {  
    // ...  
}
```



Working with the TVM Compiler

Model import



Relay Module

```
#[version = "0.0.5"]
def @main(%data : Tensor[(1, 3, 64, 64), int8],
          %weight : Tensor[(8, 3, 5, 5), int8]) {
    %1 = nn.conv2d(
        %data,
        %weight,
        padding=[2, 2],
        channels=8,
        kernel_size=[5, 5],
        data_layout="NCHW",
        kernel_layout="OIHW",
        out_dtype="int32");
    %3 = right_shift(%1, 9);
    %4 = cast(%3, dtype="int8");
    %4
}
```

Working with the TVM Compiler

Model import



Optimize Operators
(AutoTVM)

Relay Module

```
#[version = "0.0.5"]
def @main(%data : Tensor[(1, 3, 64, 64), int8],
          %weight : Tensor[(8, 3, 5, 5), int8]) {
  %1 = nn.conv2d(
    %data,
    %weight,
    padding=[2, 2],
    channels=8,
    kernel_size=[5, 5],
    data_layout="NCHW",
    kernel_layout="OIHW",
    out_dtype="int32");
  %3 = right_shift(%1, 9);
  %4 = cast(%3, dtype="int8");
  %4
}
```

TensorIR

```
primfn(placeholder_2: handle,
        placeholder_3: handle,
        T_cast_1: handle) -> ()
  allocate(kernel_vec, int8, [600]) {
    for (bs.c.fused.h.fused: int32, 0, 64)
      "parallel" {
        for (w: int32, 0, 64) {
          for (vc: int32, 0, 3) {
            data_vec[(((bs.c.fused.h.fused*192) +
              (w*3)) + vc)] =
              (uint8*)placeholder_5[(((vc*4096) +
                (bs.c.fused.h.fused*64)) + w)]
          }
        }
      }
  }
  // ...
```

Working with the TVM Compiler

Model import



Optimize Operators
(AutoTVM)

Generate C/LLVM library

Relay Module

```
#[version = "0.0.5"]
def @main(%data : Tensor[(1, 3, 64, 64), int8],
          %weight : Tensor[(8, 3, 5, 5), int8]) {
    %1 = nn.conv2d(
        %data,
        %weight,
        padding=[2, 2],
        channels=8,
        kernel_size=[5, 5],
        data_layout="NCHW",
        kernel_layout="OIHW",
        out_dtype="int32");
    %3 = right_shift(%1, 9);
    %4 = cast(%3, dtype="int8");
    %4
}
```

TensorIR


```
primfn(placeholder_2: handle,
        placeholder_3: handle,
        T_cast_1: handle) -> ()
    allocate(kernel_vec, int8, [600]) {
        for (bs.c.fused.h.fused: int32, 0, 64)
            "parallel" {
                for (w: int32, 0, 64) {
                    for (vc: int32, 0, 3) {
                        data_vec[(((bs.c.fused.h.fused*192) +
(w*3)) + vc)] =
(uint8*)placeholder_5[(((vc*4096) +
(bs.c.fused.h.fused*64)) + w)]
                    }
                }
            }
        // ...
    }
```

C Source Code

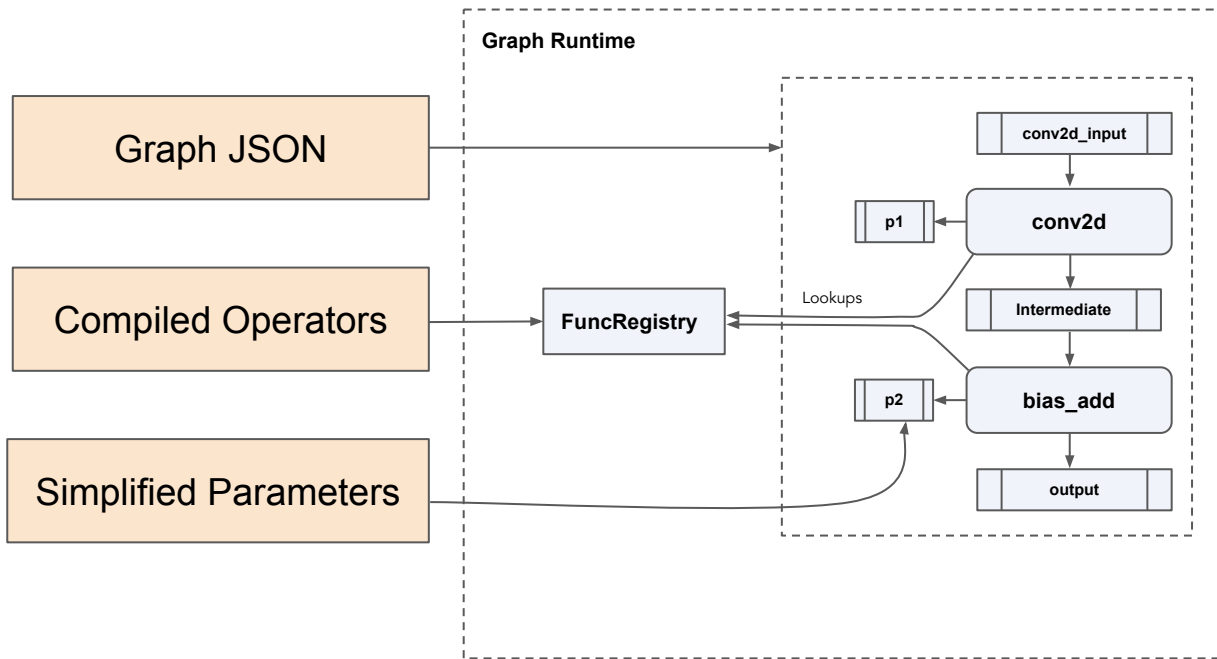
```
int32_t
fused_nn_contrib_conv2d_NCHWc_right_shift_cast(
void* args, void* arg_type_ids,
int32_t num_args, void* out_ret_value,
void* out_ret_tcode, void* resource_handle) {
    void* data_pad = TVMBackendAllocWorkspace(1,
dev_id, (uint64_t)13872, 1, 8);
    for (int32_t i0_i1_fused_i2_fused = 0;
i0_i1_fused_i2_fused < 68;
++i0_i1_fused_i2_fused) {
        for (int32_t i3 = 0; i3 < 68; ++i3) {
            for (int32_t i4 = 0; i4 < 3; ++i4) {

                ((uint8_t*)data_pad)[((((i0_i1_fused_i2_fused *
204) + (i3 * 3)) + i4))] = (((((2 <=
i0_i1_fused_i2_fused) && (i0_i1_fused_i2_fused
< 66)) && (2 <= i3)) && (i3 < 66)) ?
((uint8_t*)placeholder)[((((i0_i1_fused_i2_fus
ed * 192) + (i3 * 3)) + i4) - 390))] :
(uint8_t)0);
            }
```

Working with the TVM Compiler



Running the model end-to-end



How μ TVM Works




tvm

```
int main() {  
    // configure SoC  
    TVMInitializeRuntime();  
    TVMGraphRuntime_Run();  
}
```

+

```
int32_t fused_conv2d_right_shift_add() {  
    // ...  
}
```



Putting the pieces together

TVM Compiler Outputs

Simplified Parameters

Compiled Operators

Graph JSON

Graph Runtime

Library, from TVM

RPC Client/Server

Library, from TVM

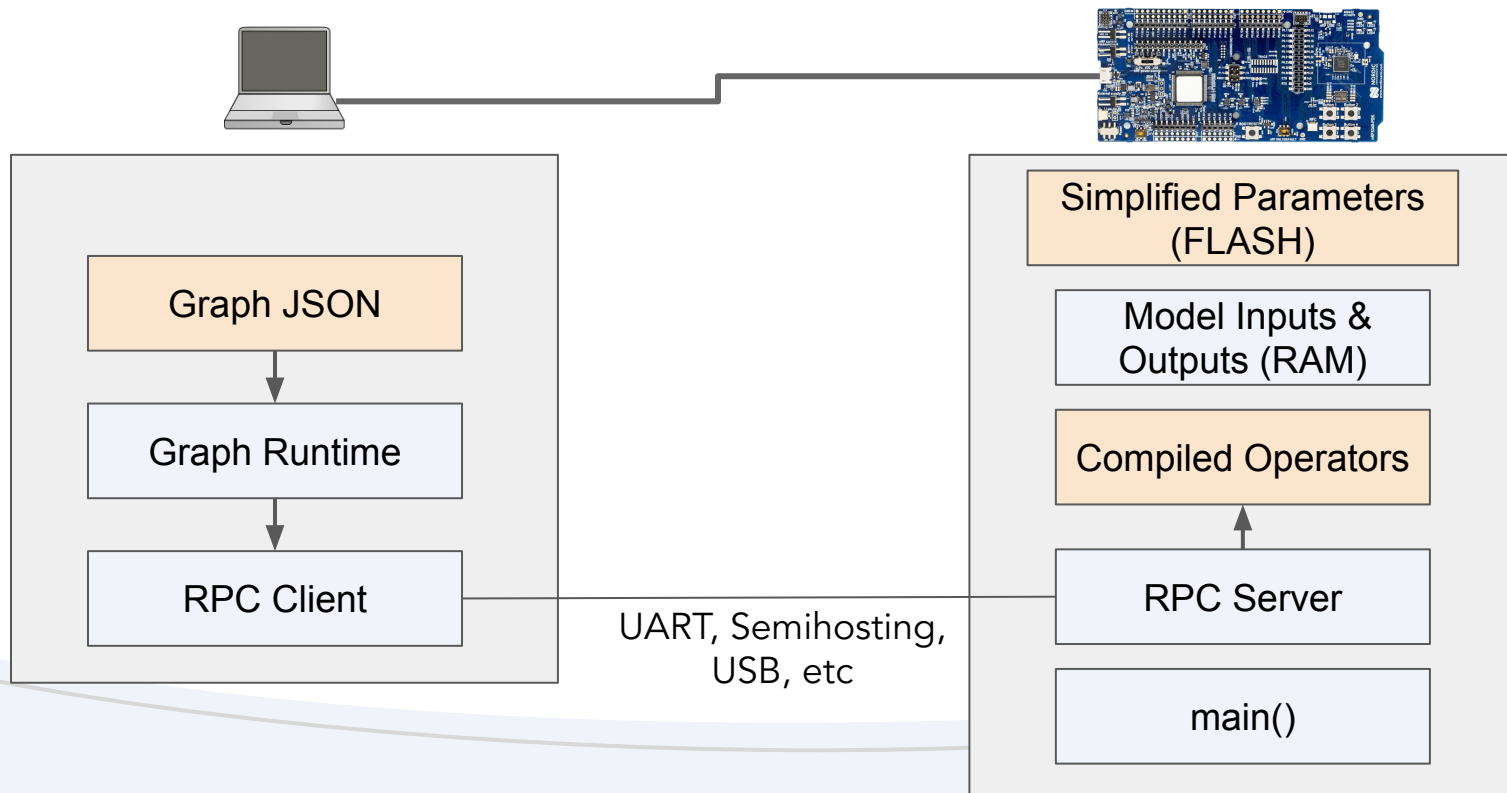
Model Inputs & Outputs (RAM)

Data, from user

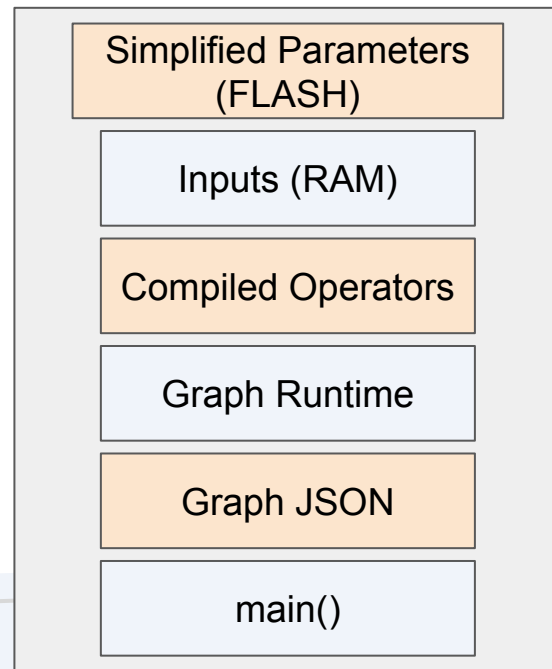

main()

SoC config & startup libraries, from RTOS/vendor

Putting the pieces together - host-driven



Putting the pieces together - standalone



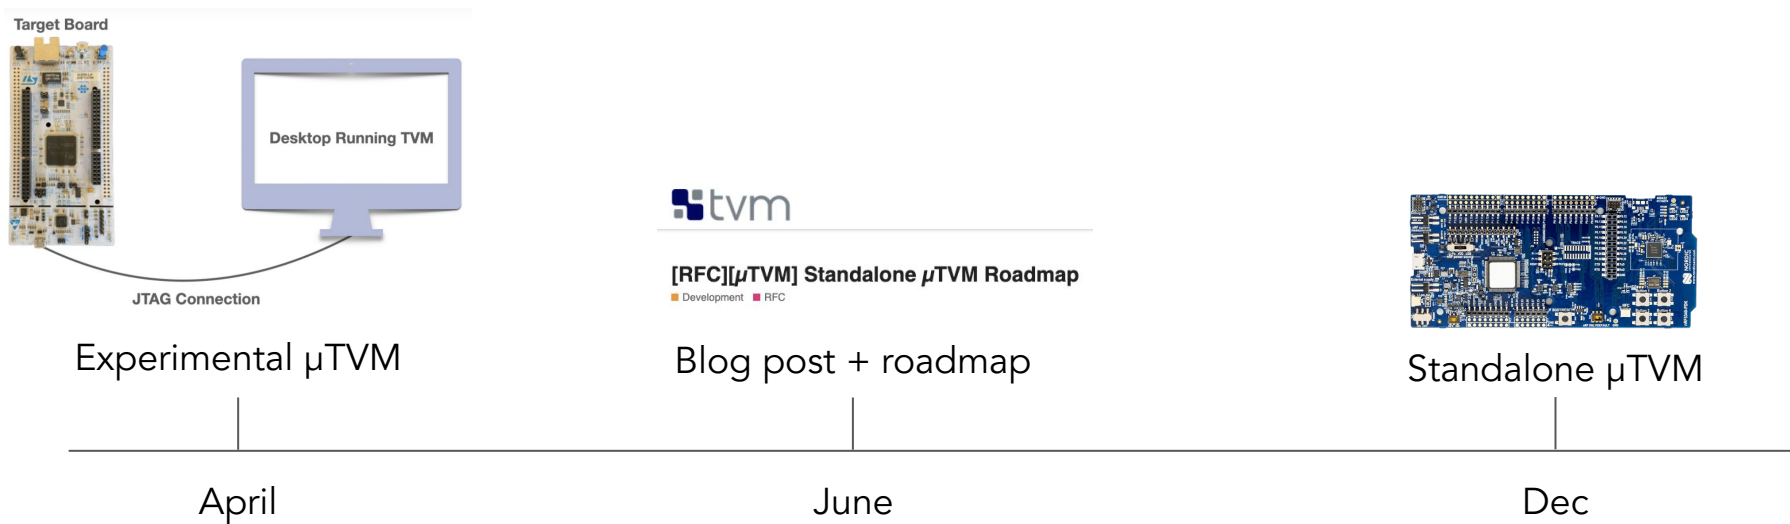
microTVM Reference Virtual Machine

- Lots of moving pieces...
 - Physical hardware
 - TVM compiler
 - GCC, LLVM, etc
 - RTOS (Zephyr, mBED), library code
 - SoC configuration / main()
- How can we collaborate?
 - Use a “Reference VM” to freeze as much of the software as possible
 - Attach hardware to VM with USB passthrough
 - See [MicroTVM Reference VM Tutorial](#) for more


Demo Walkthrough

Future Directions

μ TVM in 2020



Next for μ TVM: Ahead-of-Time Compiler



```
const DLTensor weights = {1, 2, ...};
const DLTensor biases = {4, 2, 7, ...};

int32_t classifier(DLTensor* input,
                  DLTensor* output) {
    DLTensor* intermediate =
        TVMBackendAllocWorkspace(512);


    conv2d(input, &weights, intermediate);
    bias_add(intermediate, &biases, output);

    TVMBackendFreeWorkspace(intermediate);
    return rv;
}
```

NOTE: this is a sample from the AOT compiler in development--expect API changes and an RFC.

Next for μ TVM: Hardware-aware Quantization

- Data-aware quantization v2
 - Allows quantizing more networks from within TVM
- Ultra-low-bit-width quantization
 - Could reduce the overall model memory footprint
- See [HAGO PR](#) and Ziheng Jiang's talk "Hardware-aware Quantization in TVM" on Dec 4




Riptide: Fast End-to-End Binarized Neural Networks. MLSys 2020 (March 3rd)

Joshua Fromm · Meghan Cowan · Matthai Philipose · Luis Ceze · Shwetak Patel

Next for μ TVM: Heterogeneous Execution

- Hardware acceleration offers:
 - Lower power
 - Better performance
 - More parallelism
- Potential TVM improvements:
 - CRT multi-context execution
 - TIR subgraph offloading
- See ARM's lighting talk:
"Ethos-U55 : microNPU Support for uTVM"
by Manupa Karunaratne



Next for μ TVM: Memory Planning

- Current memory planner has limitations:
 - Unaware of device memory layout
 - Requires a heap-based memory allocator
- New directions:
 - Tensor pinning
 - Accelerator-aware planning
 - Bring-your-own memory planning

Next for μ TVM: Increased Coverage

- Much of the work so far has been infrastructure-focused
- Next step: increase μ TVM coverage in terms of:
 - Supported ISA
 - Optimized Model Operators
- Auto-Scheduling can help

Next for μ TVM: Developer Experience

- Create “getting started” experience
 - Generate e.g. Arduino, Zephyr, etc projects
- Improve TVM C runtime
 - Handle faults and report through RPC server
 - Gather runtime stats to increase visibility on-device
 - Support more complex runtime scenarios -- sensing, multitasking, etc.
- Documentation
 - Targeted to developers from multiple backgrounds -- ML, firmware, etc.
 - Add design documentation and more tutorials

Q&A

- Code at <https://github.com/areusch/microtvm-blogpost-eval>
- Tutorials at <https://tvm.apache.org/docs/tutorials/index.html#micro-tvm>